

Replit board cards — how to use this set

Plain-language guide. Read this one; paste the others.

What this is

Your Replit board takes one prompt per card. These files are written to be pasted straight into "Plan a new task..." with no editing. One file equals one card.

Your board is at #13 with "Populate corrected ensemble objective results" already drafted, so these slot in after it.

Paste this one first, once

01-context-brief.md is not a task. It is shared background — the architecture rules, the current state of the database, and the things that are already verified true. Paste it as its own message to the agent at the start of a working session, or keep it open and prepend it to any card whose agent seems confused about the project.

Every card below is written to stand alone, so you do not strictly need the brief. It exists because board cards run in isolation and agents drift when they have to guess.

The important correction driving this set

Earlier I told you the GBM had no way to save predictions. That was wrong. `models/baseline.py` line 297 defines `upsert_predictions`, and `models/gbm.py` and `models/ensemble.py` both import and call it. All three models already persist to `player_prop_predictions`.

So nothing needs building. Your dashboard shows fake names because `scripts/seed_demo.py` is the only thing that has actually been run against your database. The real pipelines were never executed on real data.

That is why Card A is just "run the thing," and why it comes first. It is plausible that Card A alone fixes your empty dashboard.

Order, and why

Order	Card	What it gets you	Depends on
1	A — Run the pipeline on real data	Real names and real predicted values in the database	The 2018–2025 history load (your Done #3)
2	B — Separate demo from real	Dashboard stops showing Ava Chen and friends	A
3	C — Verify the reporting view	Proof the view behaves correctly with predictions but no odds	A, B

Order	Card	What it gets you	Depends on
4	D — Persist backtest predictions	Historical accuracy becomes visible, not just in a report file	A
5	E — Accuracy panel	Predicted vs actual on screen	D
6	F — Fix the surname alias gap	Odds name matching stops failing on compound surnames	none
7	G — Verify the OOF weight guard test	Confirms the cleanup pass actually fixed the vacuous test	your cleanup pass
8	H — ATP resumption readiness	Catches the first real odds pairing when ATP returns	F

Cards F, G, and H are independent of the dashboard work. Run them whenever.

Do A first and stop

Genuinely — run Card A, look at your dashboard, and come back. If real names appear, Cards B and C get much smaller, and you will have spent one card instead of eight. There is no point planning around a problem that a single pipeline run might resolve.

What none of these cards do

They do not attempt to backtest against historical market odds. That is not possible: prop lines for 2015–2025 were never archived and cannot be bought. Your odds history begins where your nightly capture began.

Two separate questions, worth keeping straight:

- **Accuracy** — is the prediction close to what actually happened? Answerable from historical stats alone. Cards D and E make this visible.
- **Edge** — does the prediction beat the price a book was offering? Needs the price as it stood before the match. Only collectable going forward. Card H is the readiness check for when that starts working.

A model can be accurate and still have no edge. Proving accuracy is a prerequisite, not the finish line.

Card format

Each card has: a title line you can use as the card name, a short context block, numbered tasks, explicit "do not" guardrails, and a closing instruction to report completions and omissions. That last line matters — it is the thing that has consistently produced honest reports rather than silent skips.

Note on model versions

Project rule, enforced in several cards: never overwrite an existing `modelversion`. The table has `UNIQUE (player_id, match_id, prop_type, modelversion)`, so a re-run with the same version silently updates rows in place and you lose the comparison. Bump the version on any change to features, model type, hyperparameters, or ensemble weights.